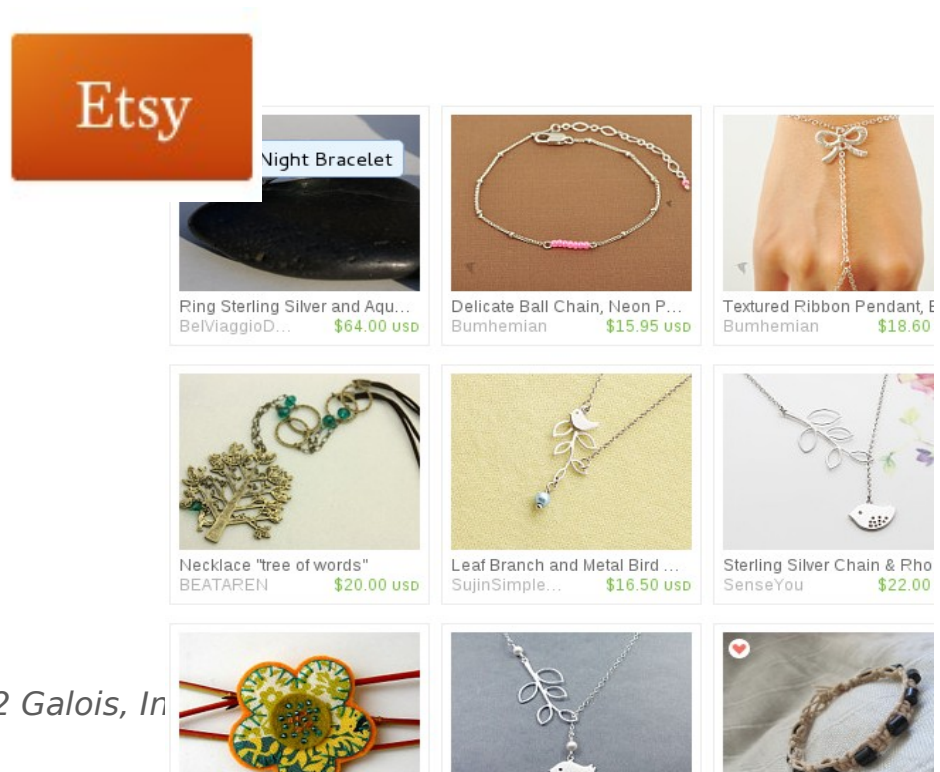


# Experience Report: a Do-it-Yourself High-Assurance Compiler

Lee Pike    Galois, Inc. <leepike@galois.com>  
Nis Wegmann    University of Copenhagen  
Sebastian Niller    Unaffiliated  
Alwyn Goodloe    NASA Langley Research Center





# Do-It-Yourself






# Do-It-Yourself

| galois |







# Do-It-Yourself



# Do-It-Yourself




Night Bracelet

Ring Sterling Silver and Aqu...  
BelViaggioD... \$64.00 USD



Delicate Ball Chain, Neon P...  
Bumhemian \$15.95 USD



Textured Ribbon Pendant  
Bumhemian \$18.60 USD




Necklace "tree of words"  
BEATAREN \$20.00 USD



Leaf Branch and Metal Bird ...  
SujinSimple... \$16.50 USD





Sterling Silver Chain & Rh...  
SenseYou \$22.00 USD





# High-Assurance

DIY



# 3 Not-So-Secret Weapons

1. Embedded domain-specific languages (EDSLs)
2. A *verifying* (not verified) compiler approach
3. Open source testing/verification libraries & tools

# National ?? and Space Administration



# National **Aeronautics** and Space Administration



# Copilot: a Run-Time Monitoring DSL

- Embedded DSL in Haskell
- Synthesize monitors for real-time embedded systems
- Stream language
- Generates Misra-like C
- Constant time, constant memory
  - Synthesized scheduler
  - No RTOS needed

# Sample Copilot specification

Haskell

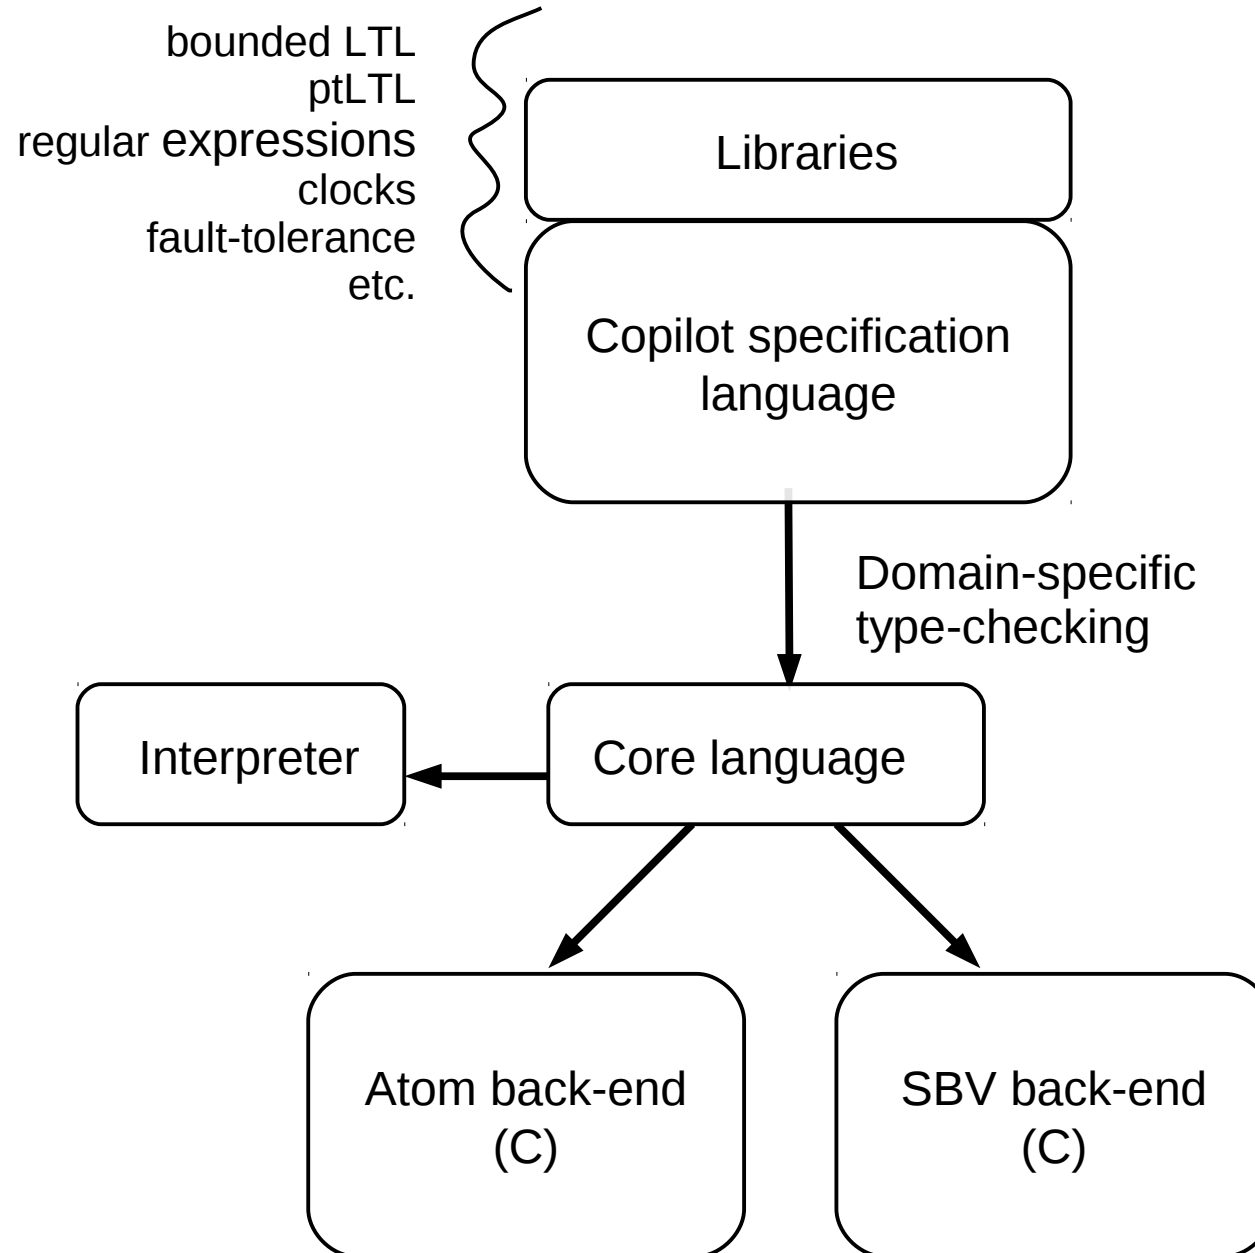
```
fib :: [Word32]
fib = [0, 1] ++ zipWith (+) fib (drop 1 fib)
```

Copilot

```
fib :: Stream Word32
fib = [0, 1] ++ (fib + drop 1 fib)
```

Special constructs for input (*sampling*) and output (*triggers*)

# Copilot Architecture





# Lessons in DIY Assurance

- Who monitors the monitor?



- Challenges:
  - EDSLs encourage rapid language design changes
  - Industrial work often doesn't “pay” for assurance (but wants it)

# Lessons in DIY Assurance


Solution: DIY assurance

- Turing *incomplete* DSLs, Turing complete macros
- Multi-level type-checking
- Cheap testing & proofs
- Unified host language

# Lesson #1: Turing-Incompleteness

Turing *incompleteness* means:

- Compiler writing is simplified
- Compiler reasoning is better (e.g., termination analysis)
- Security is improved
- Automated verification has a chance of working!




# Lesson #1: Turing-Incompleteness

Turing *incompleteness* means:

- Compiler writing is simplified
- Compiler reasoning is better (e.g., termination analysis)
- Security is improved
- Automated verification has a chance of working!

Have your cake and eat it, too:  
In an embedded DSL, the *host* language is Turing-complete!

Programs specialized at compile time.





# Lesson #2: Multi-Level Type-Checking

- Lean on Haskell's type system in the (DSL's) compiler's internal representations: e.g., GADTs
  - Leave the type system twice:
    - Pretty-print C
    - Translating between EDSLs (type-safe dynamic typing\*).
  - And ensure you aren't abusing it: **Safe Haskell**

# Lesson #2: Multi-Level Type-Checking

- Lean on Haskell's type system in the (DSL's) compiler's internal representations: e.g., GADTs
  - Leave the type system twice:
    - Pretty-print C
    - Translate between EDSLs (type-safe dynamic typing\*).
  - And ensure you aren't abusing it: **Safe Haskell**
- Then a little domain-specific type-checking:
  - Productiveness:
 

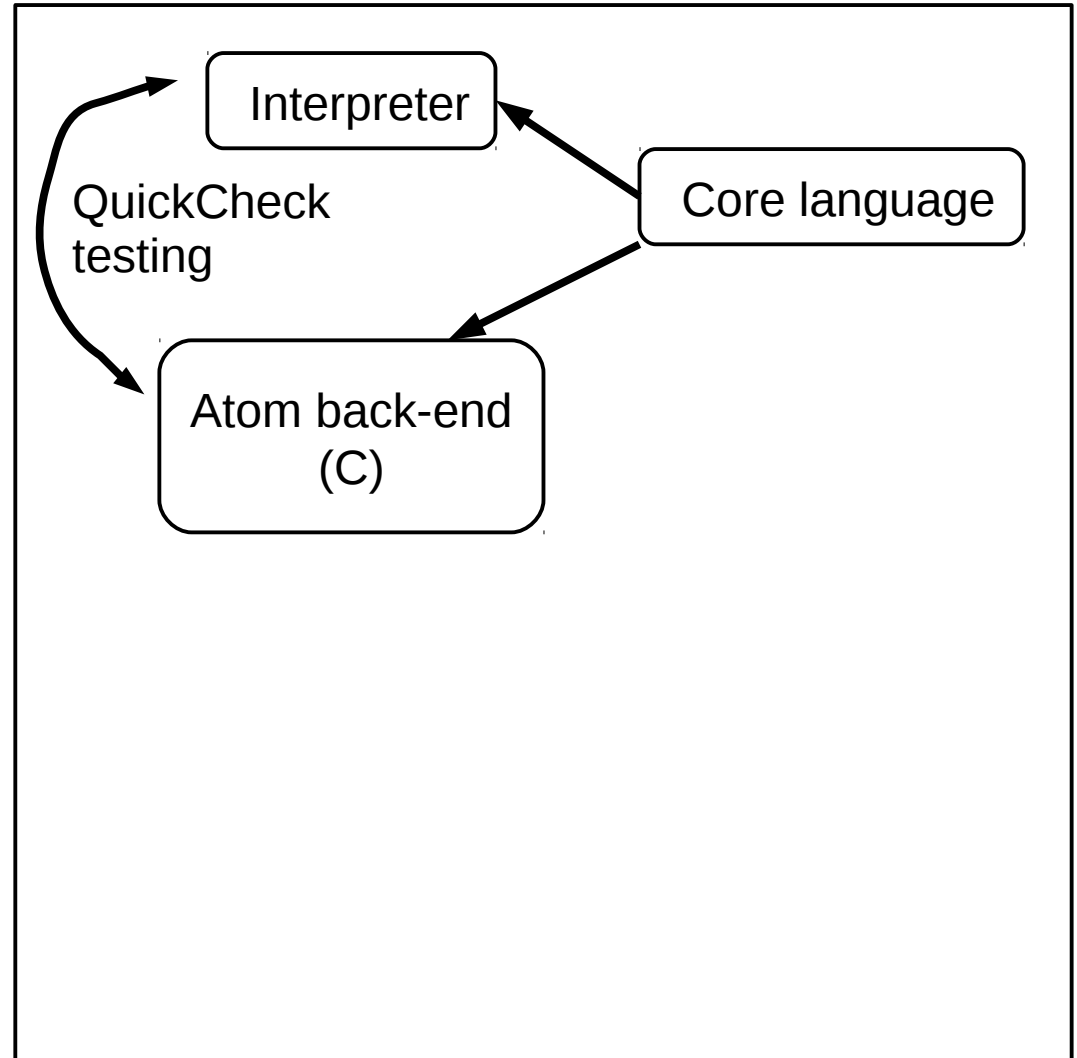
Rejected:

```
x :: Stream Word64
x = [0] ++ drop 1 x
```
  - Inputs are consistently typed (e.g., external functional calls)

# Lesson #3: Cheap Testing & Proofs

## QuickCheck:

- Small DSLs make program generation easy with good coverage
- Test ~1.5M programs/day



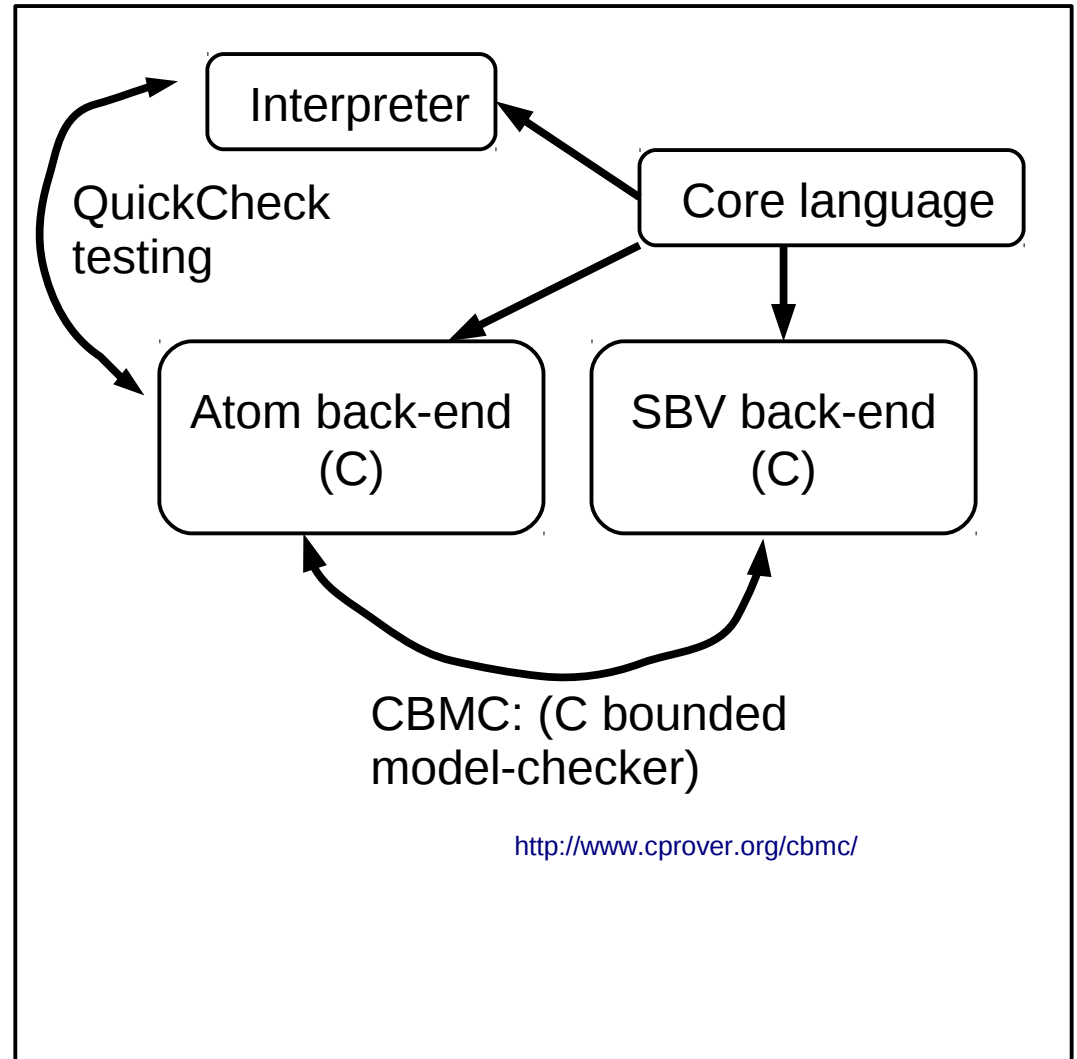
# Lesson #3: Cheap Testing & Proofs

## QuickCheck:

- Small DSLs make program generation easy with good coverage
- Test ~1.5M programs/day

## Then **prove** back-ends agree:

- Model-checking works (better) with Turing incomplete DSLs
- EDSL simplifies driver generation





# Lesson #4: a Unified Host Language

Embedded DSLs are a paradigm shift for safety-critical languages

- Fewer front-end, type-checker bugs
- “Bolting-on” new tools *within* the type system (no marshalling)
- The macro language is a build system, too!

# Conclusions



# Conclusions

## Verified compiler

- Expensive
- Specialized skills
- Hard to make repairs
- But flawless when it works



# Conclusions

## Verified compiler

- Expensive
- Specialized skills
- Hard to make repairs
- But flawless when it works



src: <http://designthatmatters.org/news/press/dtm-in-the-news/>

© 2012 Galois, Inc.



# Conclusions

## Verified compiler

- Expensive
- Specialized skills
- Hard to make repairs
- But flawless when it works



## DIY assurance

- Cheap
- Quick to build
- Easy to repair
- An “90% solution”

src: <http://designthatmatters.org/news/press/dtm-in-the-news/>



# Monitoring constraints

**Goal:** run-time monitors for software-intensive embedded systems

Runtime monitoring for real-time embedded systems should satisfy the **FaCTS**:

- **False-positives:** don't change the target's behavior
- **Certiability:** make software re-certification easy

Don't go changing sources

- **Timing:** don't interfere with the target's timing
- **SWaP:** don't exhaust size, weight, power reserves

# Software reliability is still a problem (even in ultra-critical systems)

2005-2008:

- Malaysia Airlines Flight 124 (Boeing 777)

“Software anomaly”

- Qantas Airlines Flight 72 (Airbus A330)

Transient fault in the inertial reference unit



# Software reliability is still a problem (even in ultra-critical systems)

2005-2008:

- Malaysia Airlines Flight 124 (Boeing 777)  
“Software anomaly”
- Qantas Airlines Flight 72 (Airbus A330)  
Transient fault in the inertial reference unit
- Space Shuttle STS-124 aborted launch  
Bad assumptions about distributed fault-tolerance



# Lesson #4: Safe, Unified Host Language

- Embedded DSLs are a gestalt-shift for safety-critical languages
- Few front-end bugs: (no parser, lexer, etc.)
- Type-safe translation between DSLs in the same language

E.g., seamless union of verification languages and programming languages

- The macro language is a build system, too!

```
compile program node
  (setCode (Just header)) base0pts
```

```
distCompile program node headers =
  compile (program node) node
    (setCode (Just (headers node))) base0pts
```



# 3 themes and a case-study

- RV for ultra-critical systems
  - Distributed systems
  - Hard real-time systems
  - Monitor hardware and software faults
- Using functional languages for monitor generation
  - embedded domain-specific languages (eDSL)*
- Low-cost, high assurance
- Case-study: aircraft guidance systems

# Runtime verification is needed!

How do you know your embedded software won't fail?

- Certification (e.g., DO-178B) is largely process-oriented
- Testing exercises a small fraction of the state-space
- It's probably not formally verified
  - Even if so, just a small subsystem
  - And making simplifying assumptions

**I'll argue: need the ability to detect/respond at runtime**


# Copilot Interpreter

```
evalExpr_ e0 exts locs strms = case e0 of
  Const _ x          -> x `seq` repeat x
  Drop t i id         -> strictList $
    let Just xs = lookup id strms >= fromDynF t
    in P.drop (fromIntegral i) xs
  Local t1 _ name e1 e2 -> strictList $
    let xs      = evalExpr_ e1 exts locs strms
        locs' = (name, toDynF t1 xs) : locs
    in evalExpr_ e2 exts locs' strms
  Var t name          -> strictList $
    let Just xs = lookup name locs >= fromDynF t in xs
  ExternVar t name     -> strictList $ evalExtern t name exts
  Op1 op e1             -> strictList $ repeat (evalOp1 op)
                        <*> evalExpr_ e1 exts locs strms
  Op2 op e1 e2          -> strictList $ repeat (evalOp2 op)
                        <*> evalExpr_ e1 exts locs strms
                        <*> evalExpr_ e2 exts locs strms
  Op3 op e1 e2 e3       -> strictList $ repeat (evalOp3 op)
                        <*> evalExpr_ e1 exts locs strms
                        <*> evalExpr_ e2 exts locs strms
                        <*> evalExpr_ e3 exts locs strms
```

# Flight Tests


# Experiment goals

- Monitors to check a distributed airspeed system
- Monitors also distributed & real-time  
“Bolt-on” fault-tolerance
- While satisfy timing, certifiability, SWaP goals
- Inject both physical and software faults






# Aircraft configuration Edge 540T



# Monitoring experiments


- Monitors communicate with one another over dedicated serial lines in real-time
- Properties
  - *Agreement*: return a fault-tolerant average of sensor values
    - Used to diagnose local faults
    - Diagnoses faults in the monitors **or** the sensor systems
  - Unrealistic sensor data  
Sensors values change “too fast”
- Upshot: decomposable fault-tolerance






# Monitoring results


One Byzantine-faulty processor, plus




(c) All tubes unmodified



(d) One tube stuck



(e) Two tubes stuck



(f) Three tubes stuck

# Future work

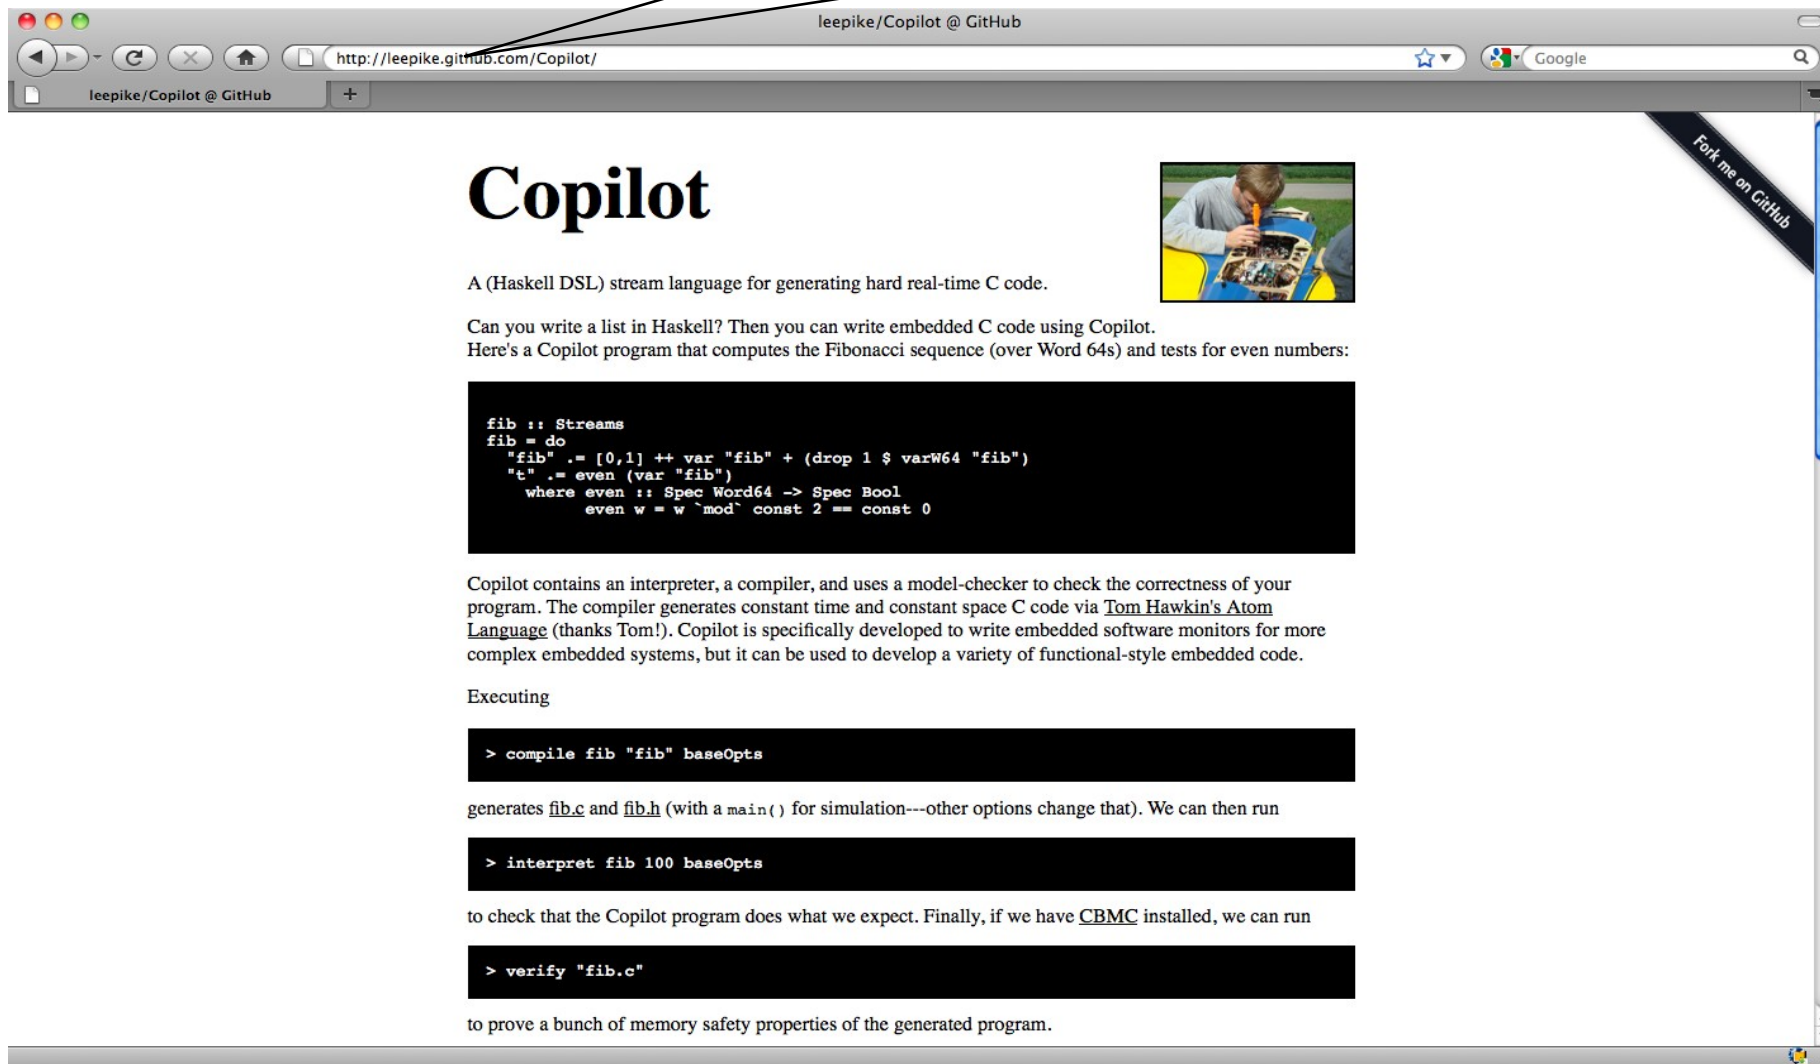
- Another case-study on autopilot communication system
- Tools for scheduling monitors
  - Used timer interrupts
  - And scheduler to decompose monitor's tasks (variable sampling, computation, etc.)
- Efficient compilation for eDSLs
- Automated mapping from real-time history to value history

E.g., state in monitor that the  $\Delta$  in  $v$  over 1sec.  $\rightarrow$  monitor maintains a history buffer of  $x$  values.

# Summary

- RV works and is needed for ultra-critical systems!
  - Distributed systems
  - Real-time systems
- Using functional languages for monitor generation
  - eDSLs*: “the benefits of functional languages applied to real-time embedded systems”
- Low-cost, high assurance

<http://leepike.github.com/Copilot/>



leepike/Copilot @ GitHub

<http://leepike.github.com/Copilot/>

# Copilot

A (Haskell DSL) stream language for generating hard real-time C code.

Can you write a list in Haskell? Then you can write embedded C code using Copilot. Here's a Copilot program that computes the Fibonacci sequence (over Word 64s) and tests for even numbers:

```
fib :: Streams
fib = do
  "fib" .= [0,1] ++ var "fib" + (drop 1 $ varW64 "fib")
  "t"  .= even (var "fib")
  where even :: Spec Word64 -> Spec Bool
        even w = w `mod` const 2 == const 0
```

Copilot contains an interpreter, a compiler, and uses a model-checker to check the correctness of your program. The compiler generates constant time and constant space C code via [Tom Hawkin's Atom Language](#) (thanks Tom!). Copilot is specifically developed to write embedded software monitors for more complex embedded systems, but it can be used to develop a variety of functional-style embedded code.

Executing

```
> compile fib "fib" baseOpts
```

generates `fib.c` and `fib.h` (with a `main()` for simulation---other options change that). We can then run

```
> interpret fib 100 baseOpts
```

to check that the Copilot program does what we expect. Finally, if we have [CBMC](#) installed, we can run

```
> verify "fib.c"
```

to prove a bunch of memory safety properties of the generated program.

Fork me on GitHub

# Differences From Lustre

- eDSL approach
- Polymorphic (embedded in Haskell)
- Simpler clock calculus—no projection operator
- BSD3
- V&V tools

# Cheap assurance

Who watches the watchmen?

- Types are free proofs—use a typed language
- Reuse existing compiler infrastructure
- Automated random testing

Ensure interpreter == compiler, millions of times

- Test coverage (line, branch, functional call) using *gcov*
- Automated back-end equivalence proofs (CBMC)

**And it's all cheap & easy.**

# Air Data Inertial Reference Units

## 35+ years of failures



### Failures cited in

- Northwest Orient Airlines Flight 6231 (1974)---3 killed
  - Increased climb/speed until uncontrollable stall
- Birgenair Flight 301, Boeing 757 (1996)---189 killed
  - One of three pitot tubes blocked; faulty air speed indicator
- Aeroperú Flight 603, Boeing 757 (1996)---70 killed
  - Tape left on the static port(!) gave erratic data
- Líneas Aèreas Flight 2553, Douglas DC-9 (1997)---74 killed
  - Freezing caused spurious low reading, compounded with a failed alarm system
  - Speed increased beyond the plane's capabilities
- Qantas Flight 72 , Airbus A330---115 injuries
  - ADIRU failure, software "limitation"
- Air France Flight 447, Airbus A330 (2009)---228 killed
  - Airspeed "unclear" to pilots
  - Still under investigation



# The power of eDSLs

- Some problems for conventional compilers go away
  - New language features are host-language macros
  - Don't need scripting languages
- E.g., compiling distributed monitors is just another host-language function:

```
compile program node
  (setCode (Just header)) baseOpts
```

```
distCompile program node headers =
  compile (program node) node
    (setCode (Just (headers node))) baseOpts
```