

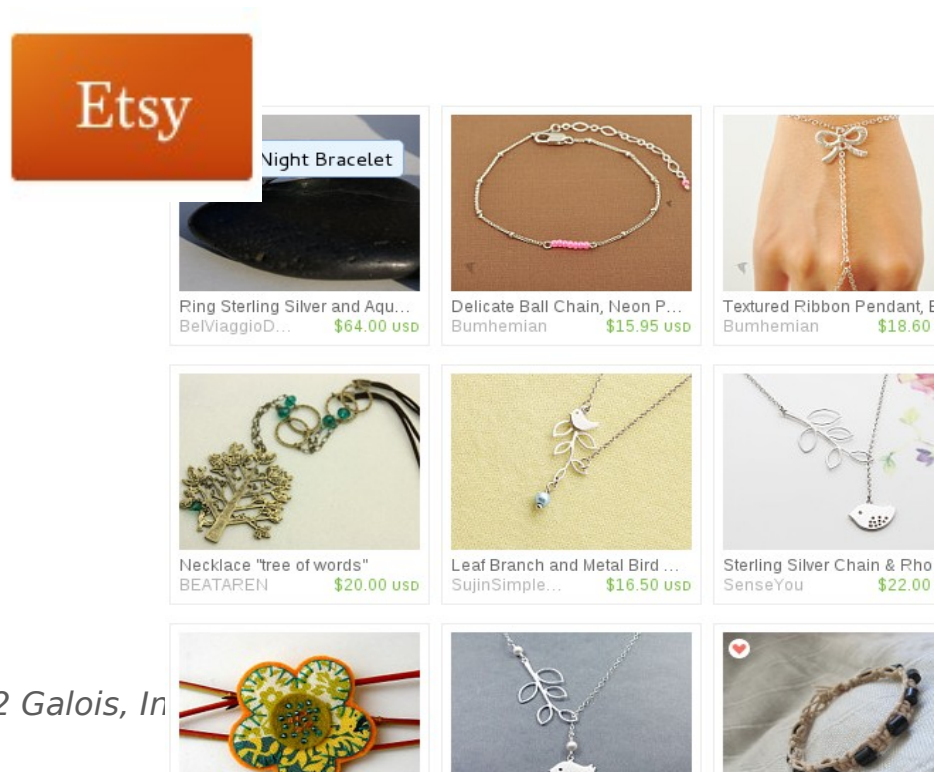
A Do-it-Yourself High-Assurance Compiler

Lee Pike Galois, Inc. <leepike@galois.com>

Nis Wegmann Technical University of Copenhagen
Sebastian Niller Unaffiliated
Alwyn Goodloe NASA Langley Research Center





Do-It-Yourself




Do-It-Yourself

| galois |



Do-It-Yourself

| galois |



DIY DRONES

Etsy


Night Bracelet



Ring Sterling Silver and Aqu...
BelViaggioD... \$64.00 USD



Delicate Ball Chain, Neon P...
Bumhemian \$15.95 USD



Textured Ribbon Pendant, B...
Bumhemian \$18.60



Necklace "tree of words"
BEATAREN \$20.00 USD




Leaf Branch and Metal Bird ...
SujinSimple... \$16.50 USD





Sterling Silver Chain & Rho...
SenseYou \$22.00



Do-It-Yourself



DIY DRONES




Night Bracelet

Ring Sterling Silver and Aqu...
BelViaggioD... \$64.00 USD



Delicate Ball Chain, Neon P...
Bumhemian \$15.95 USD



Textured Ribbon Pendant
Bumhemian \$18.60 USD




Necklace "tree of words"
BEATAREN \$20.00 USD



Leaf Branch and Metal Bird ...
SujinSimple... \$16.50 USD





Sterling Silver Chain & Rh...
SenseYou \$22.00 USD



High-Assurance

DIY



National ?? and Space Administration

National **Aeronautics** and Space Administration

Monitoring constraints

Goal: run-time monitors for software-intensive embedded systems

Runtime monitoring for real-time embedded systems should satisfy the **FaCTS**:

- **False-positives:** don't change the target's behavior
- **Certiability:** make software re-certification easy

Don't go changing sources

- **Timing:** don't interfere with the target's timing
- **SWaP:** don't exhaust size, weight, power reserves

Software reliability is still a problem (even in ultra-critical systems)

2005-2008:

- Malaysia Airlines Flight 124 (Boeing 777)
“Software anomaly”
- Qantas Airlines Flight 72 (Airbus A330)
Transient fault in the inertial reference unit
- Space Shuttle STS-124 aborted launch
Bad assumptions about distributed fault-tolerance



Our answer: Copilot

- Synthesize monitors
 - EDSL
 - From high-level specs, generate C99
Lustre-like stream language → Purely functional Misra-like C
 - Constant time: easy to compute WCET
 - Scheduler to give fine-grained timing control
 - No RTOS needed
- *Time-triggered monitoring*:
 - Sample program variables periodically
 - Keep histories as needed
 - **Not** addressing control-flow

Sample Copilot specification

Haskell

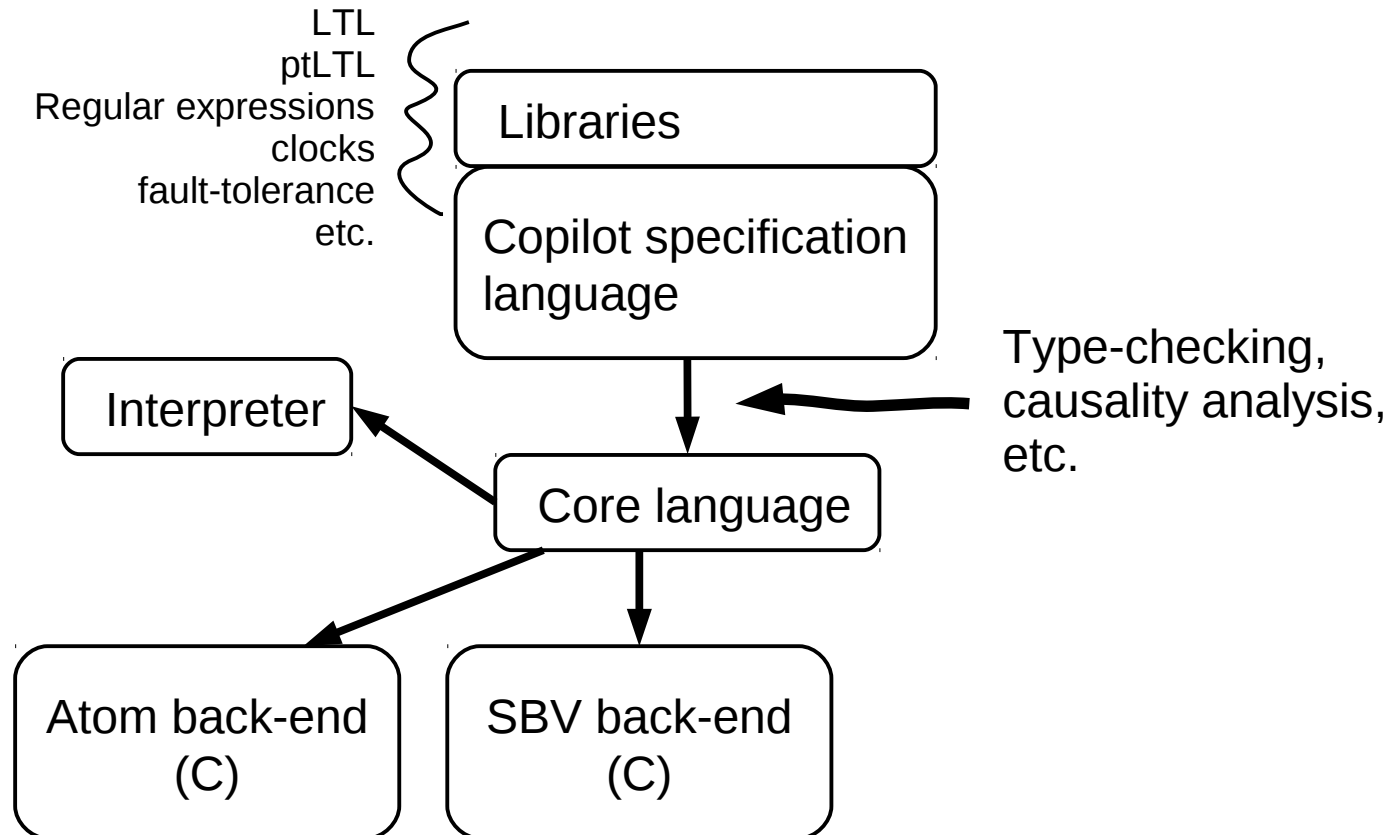
```
fib :: [Word32]
fib = [0, 1] ++ zipWith (+) fib (drop 1 fib)
```

Copilot

```
fib :: Stream Word32
fib = [0, 1] ++ (fib + drop 1 fib)
```

Special constructs for input (*sampling*) and output (*triggers*)

Copilot Architecture



Copilot: a Run-Time Monitoring DSL

- But who watches the watchmen? Copilot has to be correct
- No time for “professional” formal methods

Do-it-yourself compiler assurance?

Lessons in DIY Assurance

1. Turing *incomplete* languages, Turing complete macros
2. Multi-level type-checking
3. Cheap front-end/back-end testing
4. Unified host language

Lesson #1: Turing-Incompleteness


Turing *incompleteness* means:

- Compiler writing is simplified
- Compiler reasoning is easier (e.g., termination analysis)
- Security is improved
- Formal methods have a chance of working!

Have your cake and eat it, too:

In an embedded DSL, the *host* language is Turing-complete!

```
majority :: [Stream a] -> Stream a  
         -> Stream Word32 -> Stream a
```



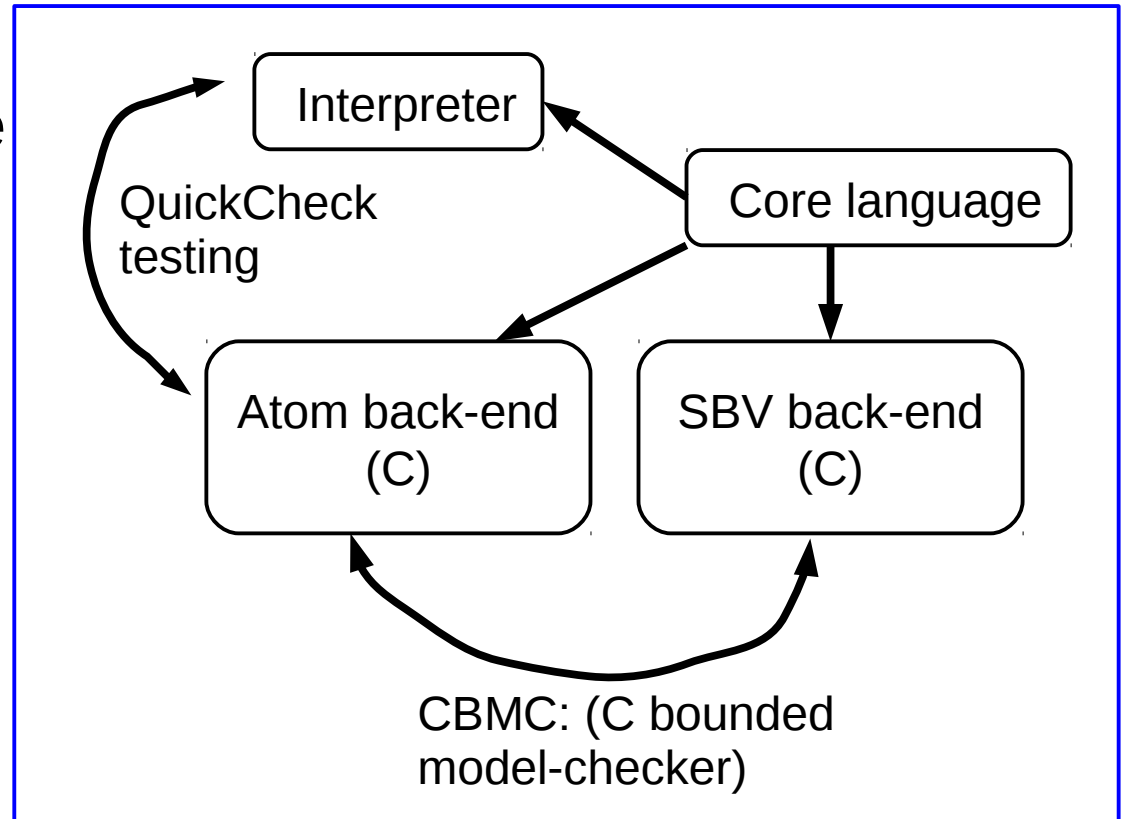
Lesson #2: Multi-Level Type-Checking

- Lean on Haskell's type system as much as possible
 - Already powerful: Polymorphism, GADTs, type-safe dynamic typing, etc.
 - But ensure you aren't abusing it: **Safe Haskell**
- Then go beyond it, e.g.:
 - Productiveness:


```
x :: Stream Word64
x = [0] ++ drop 1 x
```
 - Inputs are consistently typed (e.g., external functional calls)

Lesson #3: Cheap, High-Quality Testing

- Interpreter is the up-to-date semantics
- So make sure the back-ends agree
- Random program generation (keep the core constructs small)
- Test ~1.5M programs/day



Lesson #4: Safe, Unified Host Language

- Embedded DSLs are a gestalt-shift for safety-critical languages
- Few front-end bugs: (no parser, lexer, etc.)
- Type-safe translation between DSLs in the same language

E.g., seamless union of verification languages and programming languages

- The macro language is a build system, too!

```
compile program node
  (setCode (Just header)) base0pts
```

```
distCompile program node headers =
  compile (program node) node
    (setCode (Just (headers node))) base0pts
```

Conclusions

Compiler formal verification

- Expensive
- Specialized
- Hard to make changes
- But flawless when it works



Embedded DSLs

- Quick to prototype
- Easy to change
- Rely on existing infrastructure
- Can still be a bumpy ride...

3 themes and a case-study

- RV for ultra-critical systems
 - Distributed systems
 - Hard real-time systems
 - Monitor hardware and software faults
- Using functional languages for monitor generation
 - embedded domain-specific languages (eDSL)*
- Low-cost, high assurance
- Case-study: aircraft guidance systems

Runtime verification is needed!

How do you know your embedded software won't fail?

- Certification (e.g., DO-178B) is largely process-oriented
- Testing exercises a small fraction of the state-space
- It's probably not formally verified
 - Even if so, just a small subsystem
 - And making simplifying assumptions

I'll argue: need the ability to detect/respond at runtime


Copilot Interpreter

```
evalExpr_ e0 exts locs strms = case e0 of
  Const _ x          -> x `seq` repeat x
  Drop t i id         -> strictList $
    let Just xs = lookup id strms >= fromDynF t
    in P.drop (fromIntegral i) xs
  Local t1 _ name e1 e2 -> strictList $
    let xs      = evalExpr_ e1 exts locs strms
        locs' = (name, toDynF t1 xs) : locs
    in evalExpr_ e2 exts locs' strms
  Var t name          -> strictList $
    let Just xs = lookup name locs >= fromDynF t in xs
  ExternVar t name     -> strictList $ evalExtern t name exts
  Op1 op e1            -> strictList $ repeat (evalOp1 op)
                        <*> evalExpr_ e1 exts locs strms
  Op2 op e1 e2         -> strictList $ repeat (evalOp2 op)
                        <*> evalExpr_ e1 exts locs strms
                        <*> evalExpr_ e2 exts locs strms
  Op3 op e1 e2 e3      -> strictList $ repeat (evalOp3 op)
                        <*> evalExpr_ e1 exts locs strms
                        <*> evalExpr_ e2 exts locs strms
                        <*> evalExpr_ e3 exts locs strms
```

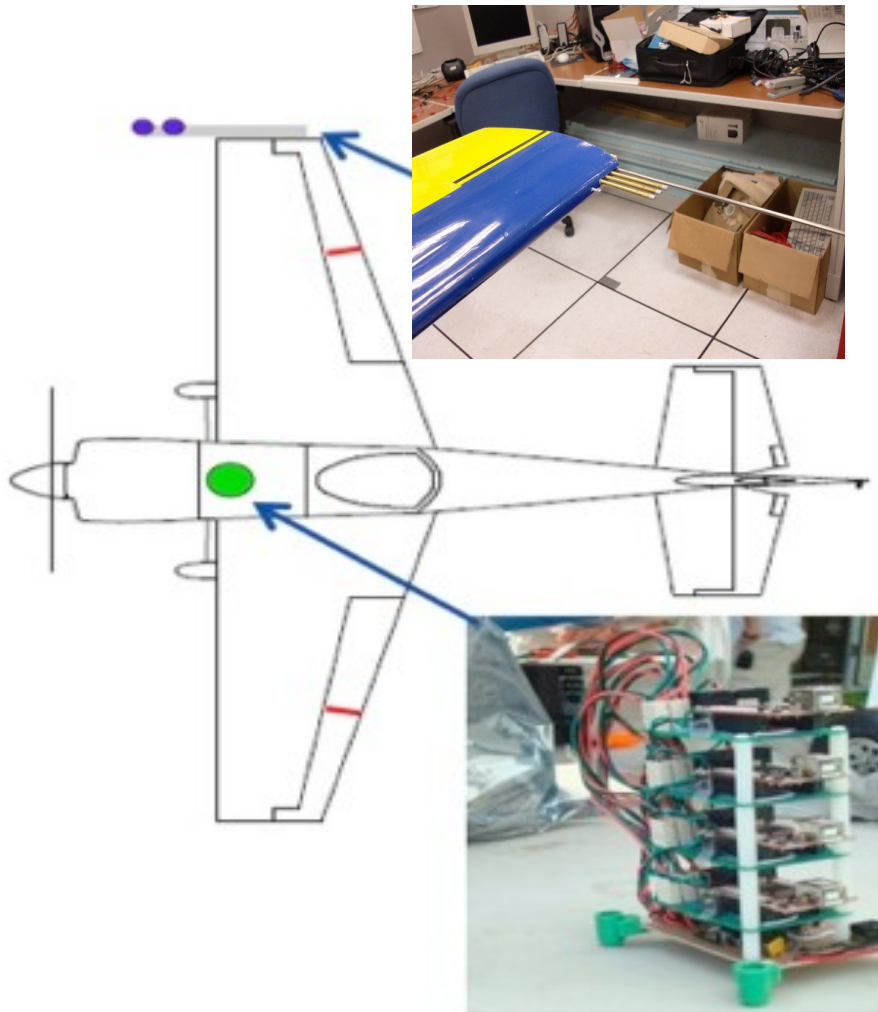

Flight Tests

Experiment goals

- Monitors to check a distributed airspeed system
- Monitors also distributed & real-time
“Bolt-on” fault-tolerance
- While satisfy timing, certifiability, SWaP goals
- Inject both physical and software faults




Aircraft configuration Edge 540T




Monitoring experiments

- Monitors communicate with one another over dedicated serial lines in real-time
- Properties
 - *Agreement*: return a fault-tolerant average of sensor values
 - Used to diagnose local faults
 - Diagnoses faults in the monitors **or** the sensor systems
 - Unrealistic sensor data
Sensors values change “too fast”
- Upshot: decomposable fault-tolerance




Monitoring results


One Byzantine-faulty processor, plus




(c) All tubes unmodified



(d) One tube stuck



(e) Two tubes stuck



(f) Three tubes stuck

Future work

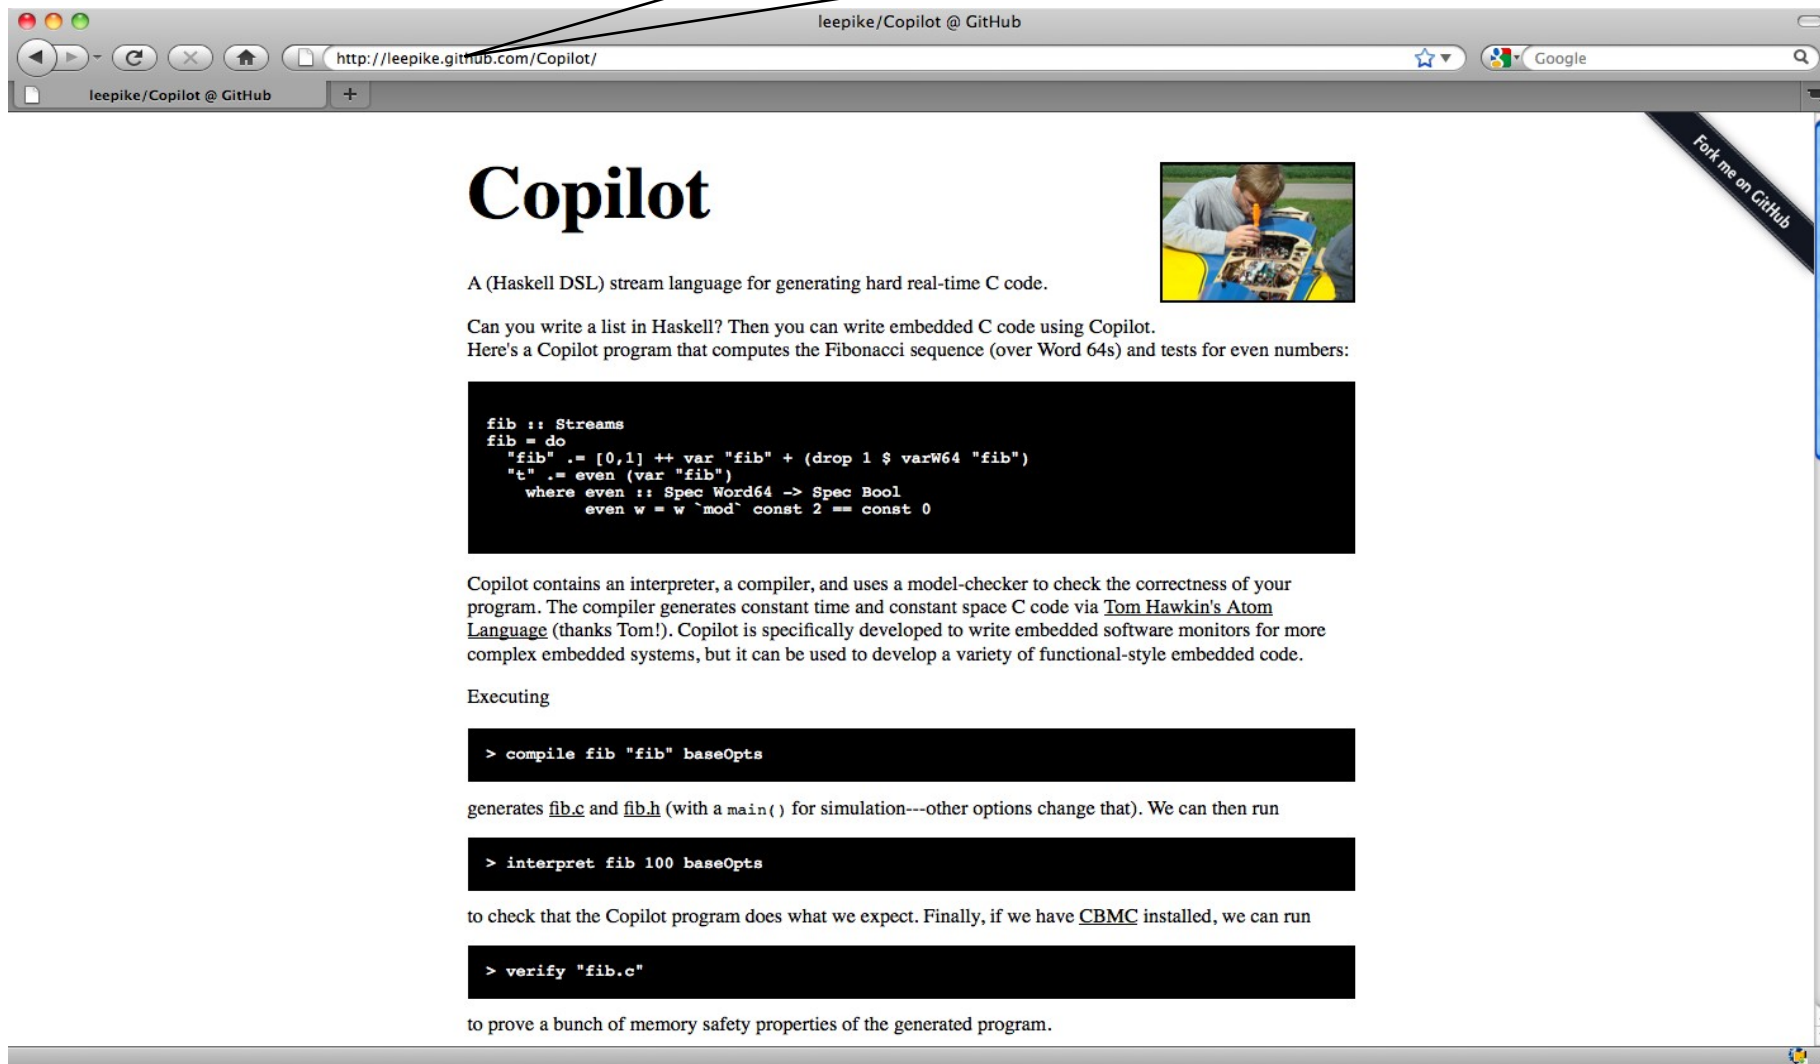
- Another case-study on autopilot communication system
- Tools for scheduling monitors
 - Used timer interrupts
 - And scheduler to decompose monitor's tasks (variable sampling, computation, etc.)
- Efficient compilation for eDSLs
- Automated mapping from real-time history to value history

E.g., state in monitor that the Δ in v over 1sec. \rightarrow monitor maintains a history buffer of x values.

Summary

- RV works and is needed for ultra-critical systems!
 - Distributed systems
 - Real-time systems
- Using functional languages for monitor generation
 - eDSLs*: “the benefits of functional languages applied to real-time embedded systems”
- Low-cost, high assurance

<http://leepike.github.com/Copilot/>



leepike/Copilot @ GitHub

<http://leepike.github.com/Copilot/>

Copilot

A (Haskell DSL) stream language for generating hard real-time C code.

Can you write a list in Haskell? Then you can write embedded C code using Copilot. Here's a Copilot program that computes the Fibonacci sequence (over Word 64s) and tests for even numbers:

```
fib :: Streams
fib = do
  "fib" .= [0,1] ++ var "fib" + (drop 1 $ varW64 "fib")
  "t"  .= even (var "fib")
  where even :: Spec Word64 -> Spec Bool
        even w = w `mod` const 2 == const 0
```

Copilot contains an interpreter, a compiler, and uses a model-checker to check the correctness of your program. The compiler generates constant time and constant space C code via [Tom Hawkin's Atom Language](#) (thanks Tom!). Copilot is specifically developed to write embedded software monitors for more complex embedded systems, but it can be used to develop a variety of functional-style embedded code.

Executing

```
> compile fib "fib" baseOpts
```

generates `fib.c` and `fib.h` (with a `main()` for simulation---other options change that). We can then run

```
> interpret fib 100 baseOpts
```

to check that the Copilot program does what we expect. Finally, if we have [CBMC](#) installed, we can run

```
> verify "fib.c"
```

to prove a bunch of memory safety properties of the generated program.

Fork me on GitHub

Differences From Lustre

- eDSL approach
- Polymorphic (embedded in Haskell)
- Simpler clock calculus—no projection operator
- BSD3
- V&V tools

Cheap assurance

Who watches the watchmen?

- Types are free proofs—use a typed language
- Reuse existing compiler infrastructure
- Automated random testing
 - Ensure interpreter == compiler, millions of times
- Test coverage (line, branch, functional call) using *gcov*
- Automated back-end equivalence proofs (CBMC)

And it's all cheap & easy.

Air Data Inertial Reference Units

35+ years of failures



Failures cited in

- Northwest Orient Airlines Flight 6231 (1974)---3 killed
 - Increased climb/speed until uncontrollable stall
- Birgenair Flight 301, Boeing 757 (1996)---189 killed
 - One of three pitot tubes blocked; faulty air speed indicator
- Aeroperú Flight 603, Boeing 757 (1996)---70 killed
 - Tape left on the static port(!) gave erratic data
- Líneas Aèreas Flight 2553, Douglas DC-9 (1997)---74 killed
 - Freezing caused spurious low reading, compounded with a failed alarm system
 - Speed increased beyond the plane's capabilities
- Qantas Flight 72 , Airbus A330---115 injuries
 - ADIRU failure, software "limitation"
- Air France Flight 447, Airbus A330 (2009)---228 killed
 - Airspeed "unclear" to pilots
 - Still under investigation

The power of eDSLs

- Some problems for conventional compilers go away
 - New language features are host-language macros
 - Don't need scripting languages
- E.g., compiling distributed monitors is just another host-language function:

```
compile program node
  (setCode (Just header)) baseOpts
```

```
distCompile program node headers =
  compile (program node) node
    (setCode (Just (headers node))) baseOpts
```